HAICHAIN

INSTITUTIONAL RISK DUE DILIGENCE REPORT V1.1 • CRYPTOGRAPHIC DEFENSE STANDARD

HAICHAIN: Institutional Risk Due Diligence Report

Mathematical Proofs, Threat Modeling & Red-Team Verification Matrix for Tier-1 Allocators and Audit Firms

TARGET AUDIENCE	CRYPTOGRAPHIC SECURITY	PROTOCOL KERNEL
Tier-1 Venture Allocators & Audit Firms	Hardware Passkey • 9 Invariants Bound	HAICHAIN v1.1 • CapLeaf 223 Bytes

HAICHAIN INSTITUTIONAL RISK DUE DILIGENCE REPORT

Mathematical Proofs, Threat Modeling & Red-Team Verification Matrix

Target Audience: Tier-1 Venture Capitalists (a16z crypto, Paradigm, Sequoia), Institutional Custodians, and Top-Tier Audit Firms (OtterSec, Neodyme, Kudelski Security)
Protocol Version: HAICHAIN Kernel v1.1 / CapLeaf v1 (223 Bytes Wire Schema)
Standard: Anchor Framework (Solana L1) + Sui Move (`AgentCap`) + Tron Rail Intent Solver

---

1. EXECUTIVE SUMMARY & SECURITY PARADIGM

Autonomous AI agents executing financial micro-transactions introduce fundamentally distinct attack surfaces from human-operated DeFi: 1. **Zero Human-in-the-Loop Latency:** An AI agent executing tool calls can be coerced via Indirect Prompt Injection into issuing thousands of unauthorized payment vouchers in seconds. 2. **Key Exfiltration Vulnerability:** Running agent instances on distributed worker nodes (Cloudflare Workers, local desktop MCPs, serverless functions) increases the risk of private key exfiltration. 3. **Sybil Counterparty Proliferation:** In an open M2M network, malicious merchants can register unlimited burner domains to bypass domain-based rate limits.

To solve this without sacrificing sub-millisecond execution, **HAICHAIN decouples execution from custody**: - **Master Authority:** WebAuthn FIDO2 Biometric Hardware Passkey (P-256) retaining 100% sovereign ownership over the spending pot. - **Sub-Agent Authority:** Ephemeral Ed25519 keypairs strictly sandboxed inside out-of-process Policy Engines with zero access to underlying wallet seeds. - **On-Chain Merkle Anchor:** 223-byte leaf wire schema enforcing 9 non-malleable mathematical invariants ($I_1 \dots I_9$).

---

2. FORMAL THREAT MODEL & LOSS BOUND FORMULAS

2.1. Attack Surface Taxonomy

Threat Class	Scenario	Conventional Web3 Exposure	HAICHAIN Worst-Case Ceiling
T_1 (Indirect Prompt Injection)	LLM coerced to pay unapproved merchant	Entire wallet drained ($100,000+$)	$E_{\text{theft}} = 0$ (Blocked by Merkle root whitelist)
T_2 (Merchant Collusion)	Legitimate merchant attempts double-claiming	Double-charge on user account	$E_{theft} = 0$ (Anchor monotonically enforces $P_i \le \min(C_i, R_i)$)
T_3 (Probation Domain Flood)	Sub-agent attempts loop drainage on new API	Cumulative exhaustion of pot	$E_{theft} \le M_{tpl}$ (Burned from parent template ceiling)
T_4 (Complete Agent Compromise)	Private key exfiltrated + colluding merchant	Total pot drained immediately	$E_{theft}(\tau) \le \min(B, K, W_p \cdot n(\tau, \omega_p))$

2.2. Theorem 6.1: Upper Bound on Worst-Case Theft (T_4)

Let τ be the total protocol response latency from the instant an unauthorized transaction is initiated to the instant the on-chain panic freeze transaction lands:

$$\tau = t_{detect} + t_{owner} + t_{land}$$

where: - $t_{detect} \leq 5.0s$ (telemetry anomaly detection threshold) - $t_{owner} \leq 2.0s$ (Passkey one-touch biometric authorization) - $t_{land} \leq 5.0s$ (Solana high-priority transaction inclusion)

Under full compromise of the agent's Ed25519 signing key and active collusion with an authenticated merchant in the active Merkle tree, the maximum cumulative theft $E_{theft}(\tau)$ is bounded by:

$$E_{theft}(\tau) \leq \min(B,\; K,\; W_p \cdot n(\tau, \omega_p),\; \Sigma_{i:\, m_i \in \mathcal{A}} E_i(\tau))$$

Where: - B: Total available pot balance (`pot.balance`). - K: Active leaf cumulative spend ceiling (`leaf.max_spend`). - W_p: Velocity limit per sliding window (`leaf.window_cap`). - ω_p: Window duration (typically $3,600s$). - $n(\tau, \omega_p) = \lceil \tau / \omega_p \rceil$: Number of sliding windows spanned during response window τ. For $\tau = 12s \ll 3600s$, $n(\tau, \omega_p) = 1$. - ΣE_i: Sum of escrow capacities locked across active channels for colluding merchants.

Corollary: For an institutional pot with $10,000$ USDT balance, a leaf budget cap of 50 USDT, and a velocity cap of 10 USDT / hour, an attacker compromising the agent's key cannot drain more than 10 USDT before the freeze takes effect, preserving $9,990$ USDT **(99.9%)** of capital.

3. RED-TEAM PENETRATION & HARDENING MATRIX

During the internal red-team forensic audit, 5 critical attack vectors were evaluated and hardened into the protocol kernel:

#	Attack Vector	Vulnerability Mechanism	Hardened On-Chain Invariant
1	Jito MEV Panic Front-Run	Malicious voucher bundles placed ahead of owner Panic Freeze	Panic Epoch Tag (<span class="m Channel tagged with epoch; instantly invalidated
2	Window Straddle Double-Dip	Splitting reserve across window boundary to draw 2x velocity	Atomic Epoch Consumption Reserve strictly burns window capacity in anchor
3	Merchant Liquidity Hostage	Uncooperative merchant leaves unspent channel un-settled	OwnerForceReclaim (<span class= Zero-claim channels unlocked immediately by master owner
4	Rent Lamport Vampire Drain	Third-party rent-recipient draining rent on channel close	Canonical Rent Recipient Rent strictly refunded to original funder (`grant_pot`)
5	Indirect Micro-Loop Sybil	1,000 child probation leaves draining small micro-reserves	Template Budget Burn `activate_child` burns parent `max_reserve` ceiling

4. MATHEMATICAL INVARIANTS AUDIT REPORT ($I_1 \dots I_9$)

Every on-chain instruction in `programs/haichain-core/src/lib.rs` preserves all 9 state invariants without exception:

1. I_1 **(Conservation of Capital):**

$$Vault_{ATA} + \Sigma_{channels}\ Escrow_j + \Sigma\ Paid_j + Refunds = Net\ Deposits$$

Verified by `test_invariant_i1_conservation` in test suite.

1. I_2 **(Canonical Payee Safety):**

$$payee_ta \equiv get_associated_token_address(merchant, mint)$$

Enforced on-chain via Anchor constraint; prevents third-party token account redirection.

1. I_3 **(Budget Ceiling):**

$$consumed \le max_spend \quad \wedge \quad budget_used \le budget_total$$

Checked before every reserve or settlement execution.

1. I_4 **(Autonomous Velocity Rate Limiter):**

$$window_used \le window_cap \quad \forall\ t \in [T_0,\ T_0 + \omega]$$

Prevents rapid automated capital drainage.

1. I_5 **(Monotonic PDA Narrowing):**

$$Channel\ PDA \subseteq Grant\ PDA \subseteq Pot\ PDA$$

Child instances cannot broaden permissions granted by parent.

1. I_6 **(Panic Epoch & Freeze Isolation):**

$$pot.panic_epoch == channel.created_at_epoch \quad \wedge \quad \neg pot.frozen$$

Instant protocol shutdown completely halts all voucher processing.

1. I_7 **(Heartbeat Expiration):**

$$current_time \le valid_until$$

Stale or abandoned sessions cannot execute new reserves.

1. I_8 **(Monotonic Settlement & Settle-Once):**

$$paid_{new} \le min(cumulative_amount,\ reserved_amount) \quad \wedge \quad paid_{new} \ge paid_{old}$$

Guarantees zero double-spending across channel vouchers.

1. I_9 **(Sybil Account Isolation):**

```
PDA("chan", grant, channel\_id) is unique
```

One capability leaf maps to exactly one on-chain account.

5. TOKENOMICS & REGULATORY COMPLIANCE

5.1. Zero-Dividend Pure Utility Standard (Howey & MiCA Immune)

- **Zero Passive Staking / Yield:** HAICHAIN has eliminated all forms of passive staking yield or dividends. The protocol offers no profit expectations derived from the entrepreneurial efforts of others.
- **50% Buyback & Burn:** 50% of protocol fees generated from network volume are used to purchase $HAISI on open DEX markets and burn it permanently, creating programmatic deflation directly tied to network utility.
- **50% Protocol Reserve Fund (CHGB):** 50% of protocol fees are directed into the autonomous slashable reserve fund to backstop provider SLAs and guarantee cross-chain liquidity.

5.2. Multi-Chain Settlement Capabilities

- **Solana L1 Anchor Kernel:** Native micro-voucher state machine with 223-byte leaf schema, running at 400ms finality.
- **Sui Move Linear Object Model (`AgentCap`):** Linear capability structs with zero `copy` or `drop` primitives, preventing duplicate delegation.
- **Tron TRC20 Rail Solver:** Off-chain intent verification with on-chain Merkle audit trail for low-cost USDT settlements in global emerging markets.

6. EXTERNAL AUDIT RECOMMENDATIONS & DATA ROOM ACCESS

The external audit package is organized into: - `/programs/haichain-core/src/lib.rs` (Anchor Rust Kernel v1.1) - `/programs/haichain-sui/sources/capleaf.move` (Sui Move Package) - `/sdk/typescript/src/capleaf.ts` (223-Byte LE Wire Schema & Merkle Tree) - `/sdk/typescript/src/policy-engine.ts` (Zero-Knowledge Key Isolation Sandbox) - `/tests/` (72 Passing Automated Fuzzing & Invariant Verification Tests)

Conclusion: HAICHAIN v1.1 provides the first mathematically proven, hardware-isolated payment layer for autonomous AI agents, setting the benchmark for institutional safety in the agentic economy.